

Drug repurposing for PCOS by drug-target interaction mining

Planning brief for the Big Data mini research proposal. Team reference, September 2026. This is the plan and the facts behind it, not the submitted text: each member writes their own proposal in their own words.

1. The assignment in one table

Section	What it must contain	Marks
Title	Disease, data, method in one line	3 (problem clarity and significance)
Biological problem	Unmet need with a number, the specific gap, the promised output. Half a page.	
Why big data	Which Vs apply, each with a size. The step that breaks a laptop.	2
Data	Real public sources, size, format, licence, one example record each	3 (data and approach)
Approach	Numbered pipeline, where Spark or distributed storage sits and why, a diagram	
Expected outcome	Concrete deliverable and a numeric success measure	2
Challenges	Hardest technical part and one honest limitation	

Length 2 to 3 pages. Hand-written or typed as instructed. Submitted through Google Classroom. The brief states that AI-generated text is detectable and receives no credit.

2. Why PCOS

Unmet need. Polycystic ovary syndrome affects roughly 1 in 10 women of reproductive age (Bozdag et al., 2016). No drug is approved specifically for PCOS. Treatment is symptomatic and borrowed: metformin (diabetes), combined oral contraceptives, spironolactone (hypertension), letrozole and clomiphene (fertility) (Teede et al., 2023). Insulin resistance, hyperandrogenism and chronic low-grade inflammation are the three biological axes (Escobar-Morreale, 2018).

The gap. The borrowed drugs were found one at a time by clinical intuition. There is no systematic, evidence-weighted ranking of all approved drugs against the PCOS gene set. That ranking is the promised output.

Why repurposing. A new drug takes 10 to 15 years and costs over a billion dollars (DiMasi et al., 2016); an approved drug already has safety and pharmacokinetic data, so a repurposing lead can move to a trial in a fraction of the time (Pushpakom et al., 2018).

The honest risk of PCOS as a choice. The disease gene set is small: about 19 to 20 GWAS loci (Day et al., 2018) and a few hundred associations in Open Targets, many with low evidence. The big-

data argument therefore cannot rest on the disease side. It rests on the drug side (all bioactivity data), the network (the full human interactome) and, optionally, expression signatures.

3. Data sources

Source	Role	Approximate size	One record	Licence
ChEMBL v34/35 (Zdrazil et al., 2023)	Drug-target bioactivities	2.4 M compounds, 20 M activities, 15 k targets. Dump ~4 GB compressed, ~25 GB loaded. Parquet on FTP.	CHEMBL25 (aspirin), target CHEMBL2094253 (PTGS1), IC50 = 32 nM, pchembl 7.5	CC BY-SA
DrugBank (Knox et al., 2023)	Approved flag, indications	16.5 k drugs (2.8 k approved small molecules), 20 k drug-target links. XML 1.5 GB.	DB00331 metformin, approved, targets PRKAB1 (activator), indication: type 2 diabetes	Academic licence
BindingDB (Liu et al., 2024)	Measured affinities, cross-check	2.9 M measurements, 1.3 M compounds, 9 k targets. TSV ~1 GB.	SMILES, UniProt P37231 (PPARG), Ki = 40 nM, PMID	CC BY 3.0
Open Targets Platform (Ochoa et al., 2022); (Buniello et al., 2024)	PCOS-gene associations with evidence scores	63 k targets, 28 k diseases, 18 M evidence strings. Parquet, ~50 GB total; associations slice ~2 GB.	ENSG00000105329 (TGFB1), EF0_0000660 (PCOS), source ot_genetics_portal, score 0.42	CC0
GWAS Catalog (Sollis et al., 2022)	Genetic seed genes	700 k associations. TSV ~300 MB.	rs11031006, FSHB, polycystic ovary syndrome, p = 1e-12	Open
STRING v12, human (Szklarczyk et al., 2022)	Protein interaction network	19.6 k proteins, 12 M scored links (1 M at score ≥ 700).	9606.ENSPO0000284981, 9606.ENSPO0000262607, combined_score 912	CC BY 4.0
LINCS L1000 (optional) (Subramanian et al., 2017)	Drug-induced expression signatures	3 M profiles, 34 k perturbagens, 978 landmark genes. Level 5 ~100+ GB.	metformin, A549, 10 μM, 978 z-scores	Open
GEO PCOS datasets (optional) (Barrett et al., 2012)	Disease expression signature	Granulosa cell and adipose microarray/RNA-seq sets, tens to hundreds of samples each.	GSE accession, sample, gene, log2 expression	Open
UniProt / Ensembl mapping (The UniProt Consortium, 2025); (Dyer et al., 2024)	Identifier bridge	Human 20 k proteins; mapping files hundreds of MB.	P37231 ↔ ENSG00000132170 ↔ 9606.ENSPO0000287820 ↔ PPARG	CC BY 4.0

Two corrections to the usual list. KEGG bulk download is licensed (Kanehisa et al., 2024), so use Reactome for pathways (Milacic et al., 2023). DisGeNET now restricts bulk access, so build the disease side on Open Targets and cite DisGeNET only as a cross-check (Piñero et al., 2019). Stating this shows the veracity check was done.


4. Why it is a big data problem, honestly

Step	Scale	Laptop?
Filter ChEMBL to approved drugs and join to targets	20 M activity rows in, ~55 k rows out	Input exceeds 8 GB RAM; needs a Spark join or a database. Output is small.
Build PCOS seed set from Open Targets + GWAS	2 GB Parquet in, ~300 to 500 genes out	Spark reads only the needed columns; result is small.
Network proximity for every approved drug on STRING	2,800 drugs × 12 M edges, degree-matched random controls (1,000 per drug)	No. ~2.8 million shortest-path computations. Embarrassingly parallel across drugs: Spark map over drug partitions with the graph broadcast.
Signature reversal against LINCS (optional)	3 M profiles, 100+ GB	No. Chunked or distributed processing required.

The four Vs, with numbers. Volume: 20 M activities, 12 M edges, optionally 3 M expression profiles. Variety: bioactivity tables, evidence strings, a graph, expression matrices, trial records, each with its own identifiers. Veracity: predicted versus experimental interactions, evidence scores from 0 to 1, licence-restricted sources excluded. Velocity: ChEMBL and Open Targets release quarterly, so the pipeline must re-run, which is an argument for automation, not for streaming.

Why Spark rather than Hadoop MapReduce. Network propagation and proximity are iterative; Spark keeps intermediate state in memory (Zaharia et al., 2016), MapReduce writes to disk each iteration (Dean & Ghemawat, 2008). Open Targets already ships Parquet, Spark's native columnar format (Apache Software Foundation, n.d.).

5. Pipeline





1. **Ingest.** Download ChEMBL (Parquet), DrugBank XML, Open Targets associations (Parquet), GWAS Catalog TSV, STRING human links. Store in HDFS or a local object store as Parquet. Tool: Spark.
2. **Approved drug set.** ChEMBL molecules with max_phase = 4, cross-checked against DrugBank "approved". Output ~2,800 drugs.
3. **Drug target sets.** ChEMBL activities with pchembl ≥ 6 (1 μ M) on human single-protein targets, joined to the approved set. Spark join over 20 M rows. Output: drug $\rightarrow$ set of UniProt accessions.
4. **PCOS seed set.** Open Targets associations for EFO_0000660 with overall score ≥ 0.1 , plus GWAS Catalog mapped genes. Weight = evidence score. Output ~300 to 500 Ensembl genes.
5. **Identifier mapping.** UniProt $\leftrightarrow$ Ensembl $\leftrightarrow$ STRING protein IDs. Log every unmapped ID; report the loss rate.
6. **Network proximity** (Guney et al., 2016). For each drug: mean shortest-path distance from its targets to the seed set on STRING (score ≥ 400), compared with 1,000 degree-matched random target sets, giving a z-score. Spark map over drugs, graph as broadcast variable.
7. **Rank and annotate.** Sort by z-score; attach for each drug its current indication, its closest PCOS genes and the evidence type behind them (the evidence path).
8. **Validate.** Before step 6, hold out drugs already used or in trials for PCOS (metformin, spironolactone, letrozole, clomiphene, pioglitazone, liraglutide, statins, inositol where listed). Measure recall of the hold-out in the top 50 and AUROC against all other approved drugs, following the trial-based validation in Cheng et al., 2018.
9. **Optional second line of evidence.** Signature reversal: PCOS granulosa or adipose signature from GEO versus LINCS drug signatures; report drugs that score on both lines.

6. Method options compared

Method	What it does	Strength	Weakness	Reference
Direct overlap (baseline)	Count of drug targets in the PCOS gene set, weighted by evidence	Trivially explainable	Only finds drugs on already-known genes	—
Network proximity (chosen)	Distance between targets and disease module versus random	Explainable, parallel, validated on many diseases	Favours well-studied proteins	(Guney et al., 2016)
Random walk with restart	Diffusion from seeds across the network	Captures multi-step effects	Iterative; parameter sensitive	(Cowen et al., 2017)
Signature reversal	Drugs whose expression signature opposes the disease signature	Genuinely big data; independent evidence line	Cell lines are not ovary or adipose	(Lamb et al., 2006); (Subramanian et al., 2017); (Sirota et al., 2011)

Knowledge-graph link prediction	Embedding model predicts drug-treats-disease edges	Powerful	Hard to explain in 3 pages; leakage risk	(Himmelstein et al., 2017)
---------------------------------	--	----------	--	----------------------------

7. Expected outcome and success measure

Deliverable. A ranked table of approved drugs with proximity z-score, current indication, closest PCOS genes and evidence type, plus a network figure for the top 10.

Success measure. At least 4 of the roughly 8 held-out drugs already used for PCOS appear in the top 50 of about 2,800, against under 0.2 expected by chance; AUROC above 0.7. If the hold-out is not recovered, the negative result still shows whether PCOS biology is reachable through the current target space, which is itself informative.

8. Challenges and limitations

- **Identifier mapping** across five sources is the hardest engineering step; a 5 to 10% loss of targets is typical and must be reported.
- **Small, noisy disease gene set.** Mitigation: weight seeds by evidence type and test whether genetic seeds alone recover the hold-out better than all seeds.
- **Literature bias.** Well-studied proteins have more edges and more evidence, so proximity favours them. Mitigation: degree-matched randomisation, as in the chosen method.
- **Data leakage.** If known PCOS drugs stay in the seed-building step, the validation is circular. Mitigation: hold-out before scoring.
- **Honest limitation.** The output is a computational prioritisation; it needs experimental and clinical validation, and cell-line expression signatures, if used, do not represent ovarian tissue.

9. Possible expansions

Direction	What it adds
LINCS signature reversal (Subramanian et al., 2017); (Sirota et al., 2011)	Second, independent evidence line; the largest data axis (100+ GB).
Evidence-type ablation (Buniello et al., 2024)	Tests whether genetic, expression or literature seeds drive the ranking; addresses veracity directly.
Side-effect filtering (SIDER, FAERS ~20 M reports) (Kuhn et al., 2015); (U.S. Food and Drug Administration, n.d.)	Down-ranks candidates with adverse events relevant to women of reproductive age.
Tissue weighting (GTEx) (The GTEx Consortium, 2020)	Weights targets by ovarian and adipose expression; cheap and biologically sensible.

Drug combinations (Cheng et al., 2019)	~4 M pairs of approved drugs; combinatorial, so genuinely distributed.
Docking of top hits (Corso et al., 2023)	Structural plausibility check for the top 10 candidates against their PCOS-relevant targets.
Literature mining (PubTator, 36 M abstracts) (Wei et al., 2024)	Published co-mentions supporting or contradicting each candidate.
Cross-disease run (PCOS and type 2 diabetes)	Shared insulin-resistance biology; turns the pipeline into a framework.
Knowledge-graph learning (Hetionet, PrimeKG) (Himmelstein et al., 2017); (Chandak et al., 2023)	The machine-learning extension, done properly.
Real-world evidence (Bycroft et al., 2018)	Incidence in cohorts of patients already on a candidate drug; access-controlled, future work.

10. Writing checklist

- Every claimed V has a size next to it.
- The laptop-breaking step is named with a memory or compute reason.
- Each source has size, format, licence and one example record.
- The identifier-mapping step is written down.
- Known PCOS drugs are held out, not reported as discoveries.
- A success number and its chance baseline are stated.
- A diagram is drawn.
- Written in your own words; the reasoning must be yours to defend.

11. References (APA 7)

Journal entries were retrieved from the DOI registry (doi.org content negotiation) on 22 September 2026; the Cheng et al. (2018) entry was completed by hand after its DOI resolved. Verify each before submission.

1. Apache Software Foundation. (n.d.). Apache Parquet documentation. <https://parquet.apache.org/docs/>
2. Barrett, T., Wilhite, S. E., Ledoux, P., Evangelista, C., Kim, I. F., Tomashevsky, M., Marshall, K. A., Phillippy, K. H., Sherman, P. M., Holko, M., Yefanov, A., Lee, H., Zhang, N., Robertson, C. L., Serova, N., Davis, S., & Soboleva, A. (2012). NCBI GEO: archive for functional genomics data sets—update. *Nucleic Acids Research*, 41(D1), D991–D995. <https://doi.org/10.1093/nar/gks1193>
3. Bozdag, G., Mumusoglu, S., Zengin, D., Karabulut, E., & Yildiz, B. O. (2016). The prevalence and phenotypic features of polycystic ovary syndrome: a systematic review and meta-analysis. *Human Reproduction*, 31(12), 2841–2855. <https://doi.org/10.1093/humrep/dew218>
4. Buniello, A., Suveges, D., Cruz-Castillo, C., Llinares, M., Cornu, H., Lopez, I., Tsukanov, K., Roldán-Romero, J., Mehta, C., Fumis, L., McNeill, G., Hayhurst, J., Martinez Osorio, R., Barkhordari, E., Ferrer, J., Carmona, M., Uniyal, P., Falaguera, M., Rusina, P., et al. (2024). Open Targets Platform: facilitating therapeutic hypotheses building in drug discovery. *Nucleic Acids Research*, 53(D1), D1467–D1475. <https://doi.org/10.1093/nar/gkae1128>

5. Bycroft, C., Freeman, C., Petkova, D., Band, G., Elliott, L. T., Sharp, K., Motyer, A., Vukcevic, D., Delaneau, O., O'Connell, J., Cortes, A., Welsh, S., Young, A., Effingham, M., McVean, G., Leslie, S., Allen, N., Donnelly, P., & Marchini, J. (2018). The UK Biobank resource with deep phenotyping and genomic data. *Nature*, 562(7726), 203–209. <https://doi.org/10.1038/s41586-018-0579-z>
6. Chandak, P., Huang, K., & Zitnik, M. (2023). Building a knowledge graph to enable precision medicine. *Scientific Data*, 10(1). <https://doi.org/10.1038/s41597-023-01960-3>
7. Cheng, F., Desai, R. J., Handy, D. E., Wang, R., Schneeweiss, S., Barabási, A.-L., & Loscalzo, J. (2018). Network-based approach to prediction and population-based validation of in silico drug repurposing. *Nature Communications*, 9(1), 2691. <https://doi.org/10.1038/s41467-018-05116-5>
8. Cheng, F., Kovács, I. A., & Barabási, A.-L. (2019). Network-based prediction of drug combinations. *Nature Communications*, 10(1). <https://doi.org/10.1038/s41467-019-09186-x>
9. Corso, G., Stärk, H., Jing, B., Barzilay, R., & Jaakkola, T. (2023). DiffDock: Diffusion steps, twists, and turns for molecular docking. *International Conference on Learning Representations (ICLR 2023)*. <https://arxiv.org/abs/2210.01189>
10. Cowen, L., Ideker, T., Raphael, B. J., & Sharan, R. (2017). Network propagation: a universal amplifier of genetic associations. *Nature Reviews Genetics*, 18(9), 551–562. <https://doi.org/10.1038/nrg.2017.38>
11. Day, F., Karaderi, T., Jones, M. R., Meun, C., He, C., Drong, A., Kraft, P., Lin, N., Huang, H., Broer, L., Magi, R., Saxena, R., Laisk, T., Urbanek, M., Hayes, M. G., Thorleifsson, G., Fernandez-Tajes, J., Mahajan, A., Mullin, B. H., et al. (2018). Large-scale genome-wide meta-analysis of polycystic ovary syndrome suggests shared genetic architecture for different diagnosis criteria. *PLOS Genetics*, 14(12), e1007813. <https://doi.org/10.1371/journal.pgen.1007813>
12. Dean, J., & Ghemawat, S. (2008). MapReduce. *Communications of the ACM*, 51(1), 107–113. <https://doi.org/10.1145/1327452.1327492>
13. DiMasi, J. A., Grabowski, H. G., & Hansen, R. W. (2016). Innovation in the pharmaceutical industry: New estimates of R&D costs. *Journal of Health Economics*, 47, 20–33. <https://doi.org/10.1016/j.jhealeco.2016.01.012>
14. Dyer, S. C., Austine-Orimoloye, O., Azov, A. G., Barba, M., Barnes, I., Barrera-Enriquez, V. P., Becker, A., Bennett, R., Beracochea, M., Berry, A., Bhai, J., Bhurji, S. K., Boddu, S., Branco Lins, P. R., Brooks, L., Ramaraju, S. B., Campbell, L. I., Martinez, M. C., Charkhchi, M., et al. (2024). Ensembl 2025. *Nucleic Acids Research*, 53(D1), D948–D957. <https://doi.org/10.1093/nar/gkae1071>
15. Escobar-Morreale, H. F. (2018). Polycystic ovary syndrome: definition, aetiology, diagnosis and treatment. *Nature Reviews Endocrinology*, 14(5), 270–284. <https://doi.org/10.1038/nrendo.2018.24>
16. Guney, E., Menche, J., Vidal, M., & Barabási, A.-L. (2016). Network-based in silico drug efficacy screening. *Nature Communications*, 7(1). <https://doi.org/10.1038/ncomms10331>
17. Himmelstein, D. S., Lizée, A., Hessler, C., Brueggeman, L., Chen, S. L., Hadley, D., Green, A., Khankhanian, P., & Baranzini, S. E. (2017). Systematic integration of biomedical knowledge prioritizes drugs for repurposing. *eLife*, 6. CLOCKSS. <https://doi.org/10.7554/elife.26726>
18. Kanehisa, M., Furumichi, M., Sato, Y., Matsuura, Y., & Ishiguro-Watanabe, M. (2024). KEGG: biological systems database as a model of the real world. *Nucleic Acids Research*, 53(D1), D672–D677. <https://doi.org/10.1093/nar/gkae909>
19. Knox, C., Wilson, M., Klinger, C., Franklin, M., Oler, E., Wilson, A., Pon, A., Cox, J., Chin, N. E., Strawbridge, S., Garcia-Patino, M., Kruger, R., Sivakumaran, A., Sanford, S., Doshi, R., Khetarpal, N., Fatokun, O., Doucet, D., Zubkowski, A., et al. (2023). DrugBank 6.0: the DrugBank Knowledgebase for 2024. *Nucleic Acids Research*, 52(D1), D1265–D1275. <https://doi.org/10.1093/nar/gkad976>
20. Kuhn, M., Letunic, I., Jensen, L. J., & Bork, P. (2015). The SIDER database of drugs and side effects. *Nucleic Acids Research*, 44(D1), D1075–D1079. <https://doi.org/10.1093/nar/gkv1075>
21. Lamb, J., Crawford, E. D., Peck, D., Modell, J. W., Blat, I. C., Wrobel, M. J., Lerner, J., Brunet, J.-P., Subramanian, A., Ross, K. N., Reich, M., Hieronymus, H., Wei, G., Armstrong, S. A., Haggarty, S. J., Clemons, P. A., Wei, R., Carr, S. A., Lander, E. S., & Golub, T. R. (2006). The Connectivity Map: Using Gene-Expression Signatures to Connect Small Molecules, Genes, and Disease. *Science*, 313(5795), 1929–1935. <https://doi.org/10.1126/science.1132939>
22. Liu, T., Hwang, L., Burley, S., Nitsche, C., Southan, C., Walters, W., & Gilson, M. (2024). BindingDB in 2024: a FAIR knowledgebase of protein-small molecule binding data. *Nucleic Acids Research*, 53(D1), D1633–D1644. <https://doi.org/10.1093/nar/gkae1075>

23. Milacic, M., Beavers, D., Conley, P., Gong, C., Gillespie, M., Griss, J., Haw, R., Jassal, B., Matthews, L., May, B., Petryszak, R., Ragueneau, E., Rothfels, K., Sevilla, C., Shamovsky, V., Stephan, R., Tiwari, K., Varusai, T., Weiser, J., et al. (2023). The Reactome Pathway Knowledgebase 2024. *Nucleic Acids Research*, 52(D1), D672–D678. <https://doi.org/10.1093/nar/gkad1025>
24. National Library of Medicine. (n.d.). ClinicalTrials.gov [Database]. <https://clinicaltrials.gov>
25. Ochoa, D., Hercules, A., Carmona, M., Suveges, D., Baker, J., Malangone, C., Lopez, I., Miranda, A., Cruz-Castillo, C., Fumis, L., Bernal-Llinares, M., Tsukanov, K., Cornu, H., Tsigos, K., Razuvayevskaya, O., Buniello, A., Schwartzentruber, J., Karim, M., Ariano, B., et al. (2022). The next-generation Open Targets Platform: reimaged, redesigned, rebuilt. *Nucleic Acids Research*, 51(D1), D1353–D1359. <https://doi.org/10.1093/nar/gkac1046>
26. Piñero, J., Ramírez-Anguita, J. M., Saüch-Pitarch, J., Ronzano, F., Centeno, E., Sanz, F., & Furlong, L. I. (2019). The DisGeNET knowledge platform for disease genomics: 2019 update. *Nucleic Acids Research*. <https://doi.org/10.1093/nar/gkz1021>
27. Pushpakom, S., Iorio, F., Eyers, P. A., Escott, K. J., Hopper, S., Wells, A., Doig, A., Guilleams, T., Latimer, J., McNamee, C., Norris, A., Sanseau, P., Cavalla, D., & Pirmohamed, M. (2018). Drug repurposing: progress, challenges and recommendations. *Nature Reviews Drug Discovery*, 18(1), 41–58. <https://doi.org/10.1038/nrd.2018.168>
28. Sirota, M., Dudley, J. T., Kim, J., Chiang, A. P., Morgan, A. A., Sweet-Cordero, A., Sage, J., & Butte, A. J. (2011). Discovery and Preclinical Validation of Drug Indications Using Compendia of Public Gene Expression Data. *Science Translational Medicine*, 3(96). <https://doi.org/10.1126/scitranslmed.3001318>
29. Sollis, E., Mosaku, A., Abid, A., Buniello, A., Cerezo, M., Gil, L., Groza, T., Güneş, O., Hall, P., Hayhurst, J., Ibrahim, A., Ji, Y., John, S., Lewis, E., MacArthur, J. L., McMahon, A., Osumi-Sutherland, D., Panoutsopoulou, K., Pendlington, Z., et al. (2022). The NHGRI-EBI GWAS Catalog: knowledgebase and deposition resource. *Nucleic Acids Research*, 51(D1), D977–D985. <https://doi.org/10.1093/nar/gkac1010>
30. Subramanian, A., Narayan, R., Corsello, S. M., Peck, D. D., Natoli, T. E., Lu, X., Gould, J., Davis, J. F., Tubelli, A. A., Asiedu, J. K., Lahr, D. L., Hirschman, J. E., Liu, Z., Donahue, M., Julian, B., Khan, M., Wadden, D., Smith, I. C., Lam, D., et al. (2017). A Next Generation Connectivity Map: L1000 Platform and the First 1,000,000 Profiles. *Cell*, 171(6), 1437–1452. <https://doi.org/10.1016/j.cell.2017.10.049>
31. Szklarczyk, D., Kirsch, R., Koutrouli, M., Nastou, K., Mehryary, F., Hachilif, R., Gable, A. L., Fang, T., Doncheva, N., Pyysalo, S., Bork, P., Jensen, L., & von Mering, C. (2022). The STRING database in 2023: protein–protein association networks and functional enrichment analyses for any sequenced genome of interest. *Nucleic Acids Research*, 51(D1), D638–D646. <https://doi.org/10.1093/nar/gkac1000>
32. Teede, H. J., Tay, C. T., Laven, J. J. E., Dokras, A., Moran, L. J., Piltonen, T. T., Costello, M. F., Boivin, J., Redman, L. M., Boyle, J. A., Norman, R. J., Mousa, A., Joham, A. E., Arlt, W., Azziz, R., Balen, A., Bedson, L., Berry, L., et al. (2023). Recommendations from the 2023 international evidence-based guideline for the assessment and management of polycystic ovary syndrome. *European Journal of Endocrinology*, 189(2), G43–G64. <https://doi.org/10.1093/ejendo/lvad096>
33. The GTEx Consortium. (2020). The GTEx Consortium atlas of genetic regulatory effects across human tissues. *Science*, 369(6509), 1318–1330. <https://doi.org/10.1126/science.aaz1776>
34. The UniProt Consortium. (2025). UniProt: The Universal Protein Knowledgebase in 2025. *Nucleic Acids Research*, 53(D1), D609–D617. <https://doi.org/10.1093/nar/gkae1010>
35. U.S. Food and Drug Administration. (n.d.). FDA Adverse Event Reporting System (FAERS) public dashboard [Database]. <https://fis.fda.gov/>
36. Wei, C.-H., Allot, A., Lai, P.-T., Leaman, R., Tian, S., Luo, L., Jin, Q., Wang, Z., Chen, Q., & Lu, Z. (2024). PubTator 3.0: an AI-powered literature resource for unlocking biomedical knowledge. *Nucleic Acids Research*, 52(W1), W540–W546. <https://doi.org/10.1093/nar/gkae235>
37. Zaharia, M., Xin, R. S., Wendell, P., Das, T., Armbrust, M., Dave, A., Meng, X., Rosen, J., Venkataraman, S., Franklin, M. J., Ghodsi, A., Gonzalez, J., Shenker, S., & Stoica, I. (2016). Apache Spark. *Communications of the ACM*, 59(11), 56–65. <https://doi.org/10.1145/2934664>
38. Zdrazil, B., Felix, E., Hunter, F., Manners, E. J., Blackshaw, J., Corbett, S., de Veij, M., Ioannidis, H., Lopez, D. M., Mosquera, J., Magarinos, M., Bosc, N., Arcila, R., Kizilören, T., Gaulton, A., Bento, A., Adasme, M., Monecke, P., Landrum, G., & Leach, A. (2023). The ChEMBL Database in 2023: a drug discovery platform spanning multiple bioactivity data types and time periods. *Nucleic Acids Research*, 52(D1), D1180–D1192. <https://doi.org/10.1093/nar/gkad1004>

Dataset sizes are approximate, from 2025 releases. Cite the release you download.